

RNA-Seq with Bioconductor in R

by Mary Piper and DataCamp


The Genome

all living organisms contain the instructions for life in their genome
which is present in the nuclei of their cells
the genome is comprised of DNA divided into chromosomes
the building blocks of DNA are called nucleotides
our bases are guanine, adenine, cytosine, and thymine
DNA is double-stranded
forms a helix with a sugar-phosphate backbone
A to T
G to C
the order of these pairs is referred to as the DNA sequence

Genes

with the DNA sequence are regions called genes
genes provide instructions to make proteins
proteins perform some function within the cell
to make proteins, DNA is transcribed into messenger RNA (mRNA)
this is translated by the ribosome into protein
some genes encode RNA that does not get translated into protein called non-coding RNAs (ncRNA)
ncRNAs have a function in and of themselves and include rRNAs, tRNAs, siRNAs, and others
All RNAs transcribed from genes are called transcripts

RNA processing



Wikimedia Commons: [Gene structure eukaryote 2 annotated.svg](#) by Thomas Shafee, used under Creative Commons Attribution 4.0 International

to be translated into proteins, mRNA must undergo processing
above describes the process

the top strand represents a gene in the DNA

UTR represents the untranslated regions

after post-transcriptional processing, the introns are spliced out and a polyA tail
5'cap are added to yield mature mRNA transcripts

these mature mRNA transcripts can be translated into protein

mRNA transcripts have a polyA tail, which is a sequence of AAAs at the end stating
the end of the transcript

many non-coding RNAs do not have this

we have different cells due to the different genes that are turned on in these cells

Gene expression in disease

a mutation can affect the type and quantity of RNAs and proteins produced

*to explore gene expression changes that occur in disease, it can be useful to
measure the quantity of RNA expressed by all genes

we can do this by using RNA-Seq

then differentail expression analysis of RNA-Seq data can be used to determine
whether there are significant differences in gene expression between conditions

RNA-Seq questions

what genes are differentially expressed between sample groups?

are there any trends in gene expression over time or across conditions?

which groups of genes change similarly over time or across conditions?

what processes or pathways are enriched for my condition of interest?

Need libraries for RNA-Seq workflow

Load library for DESeq2

library(DESeq2)

Load library for RColorBrewer

library(RColorBrewer)

Load library for pheatmap

library(pheatmap)

Load library for tidyverse

library(tidyverse)

RNA-Seq Workflow

Experimental Design

'Technical replicates' are of low variation

*'Biological replicates' are key, crucial to the success of analyses

the more the better but at the very least have to have 3

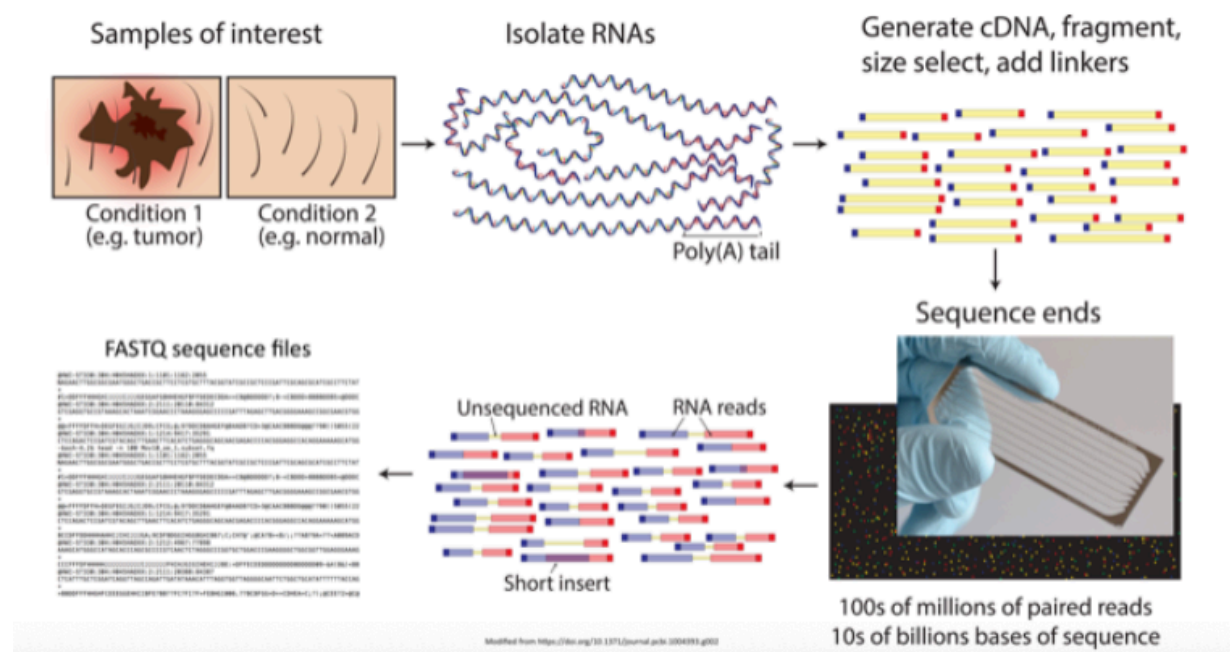
*try to avoid experiments that are performed batch-wise

goal is to perform experiment across all conditions at the same time

if cannot avoid distribute samples from each sample group into each batch

avoid major sources of variation (ie mix sexes)

Workflow



RNA isolated and DNA contamination removed

rRNA removed or mature mRNAs are selected by their polyA tails
 RNA is turned into cDNA > fragmented, size selected, and adapters are added
 done to generate the RNA-Seq libraries to be sequenced
 sequencing generates millions of nucleotide sequences called reads
 reads correspond to ends of the fragments sequenced
 the sequence of each read is output into fastq files
 now we assess the quality of the raw data
 ensuring nothing went wrong at the sequencing facility
 looking for data contamination
 next is alignment or mapping of the reads to the genome to determine the location
 on the genome where the reads originated

More on alignment

mRNA contains only the exons needed to create the proteins
 when mRNA is aligned to the genome containing introns, reads will be split across
 introns
 tools for aligning reads to the genome need to align across introns for RNA-seq
 output of alignment gives the genome coordinates for where the read most likely
 originated from in the genome

Next the reads aligning to the exons of each gene are quantified to yield a matrix of gene counts

**up to this point we need to use command line tools
 from here we can move into using R and Bioconductor
 example

`we_rawcounts <- read.csv("fibrosis_wt_rawcounts.csv")`

this example output > samples as columns and gene IDs as rows

	wt_normal1	wt_normal2	wt_normal3	wt_fibrosis1	wt_fibrosis2	wt_fibrosis3	wt_fibrosis4
ENSMUSG00000102693	0	0	0	0	0	0	0
ENSMUSG00000064842	0	0	0	0	0	0	0
ENSMUSG00000051951	3	1	1	42	52	16	35
ENSMUSG00000102851	0	0	0	0	0	0	0
ENSMUSG00000103377	0	0	0	0	0	0	0

here the count values represent the number of reads or fragments aligning to the exons of each gene

Example of statistical analysis

	baseMean	log2FoldChange	lfcSE	stat	pvalue	padj
MOV10	21681.7998	4.7695983	0.10269615	46.232357	0.000000e+00	0.000000e+00
H1FO	7881.0811	1.5250811	0.05548216	27.479961	3.047848e-166	2.489330e-162
HIST1H1C	1741.3830	1.4868361	0.06844630	21.700664	2.022230e-104	1.101104e-100
TXNIP	5133.7486	1.3868320	0.06759178	20.513587	1.628305e-93	6.649590e-90
NEAT1	21973.7061	0.9087853	0.04601897	19.747620	8.408861e-87	2.747175e-83
KLF10	1694.2109	1.2093969	0.06339756	19.067600	4.693529e-81	1.277813e-77
INSIG1	11872.5106	1.2260848	0.06780306	18.079993	4.581384e-73	1.069099e-69

includes log2 fold changes of expression between conditions and the adjusted p-values for each gene

genes that reach a threshold for significance can be subset to define a list of significant differentially expressed genes of potential interest

More on fold change

A fold change is a measure that describes how much a quantity changes between two conditions. In the context of genomics and gene expression studies, a fold change is often used to quantify the difference in expression levels of a gene under two different experimental conditions, such as a treatment group compared to a control group.

The formula for calculating fold change is straightforward:

$$\text{Fold Change} = \frac{\text{Expression in Condition A}}{\text{Expression in Condition B}}$$

Here's how to interpret fold change:

- If the fold change is 1, it indicates no change in expression between the two conditions.
- If the fold change is greater than 1, it suggests upregulation or an increase in expression in the first condition compared to the second.
- If the fold change is less than 1, it suggests downregulation or a decrease in expression in the first condition compared to the second.

For example, a fold change of 2 means that the expression level in Condition A is two times higher than in Condition B. Similarly, a fold change of 0.5 means that the expression level in Condition A is half that of Condition B.

Fold change is a valuable metric in genomics because it provides a relative measure of the magnitude of differences in gene expression between experimental conditions. Researchers often use fold changes along with statistical tests to

identify genes that show significant changes in expression under specific conditions.

Example

```
# Explore the first six observations of smoc2_rawcounts  
head(smoc2_rawcounts)
```

```
# Explore the structure of smoc2_rawcounts  
str(smoc2_rawcounts)
```

Differential gene expression overview

we use DESeq2 to determine whether genes are expressed at significantly different levels between two or more sample groups

the question off the bat is always why not just identify the genes with the largest fold changes in expression between sample groups?

the answer is because we need to account for variation

*goal is to determine for each gene whether the differences in expression between groups is significant given the amount of variation within groups (or between the biological replicates)

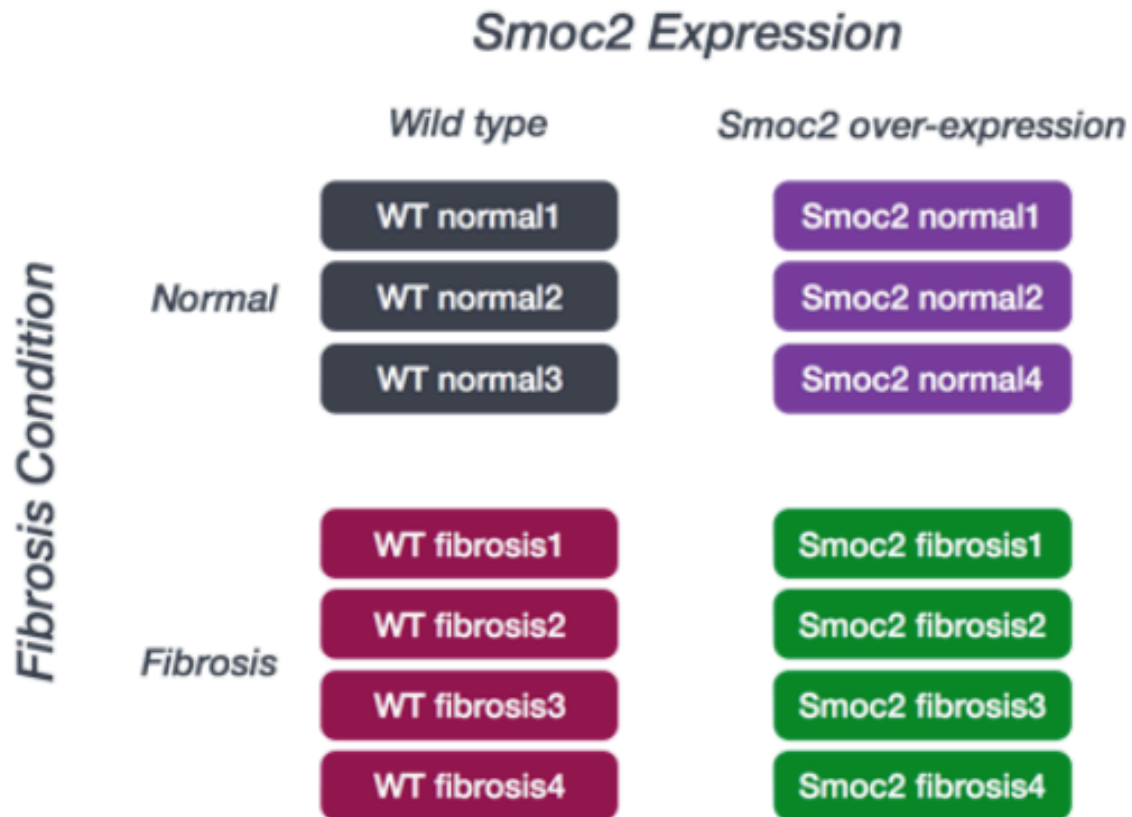
Example that we are going to work on is the Smoc2 gene

ie Secreted modular calcium-binding protein 2

noted to have increased expression in kidney fibrosis

kidney fibrosis is characterized by an excess of extracellular matrix in the space between tubules and capillaries

Our dataset



Picking the appropriate statistical model > **this is determined by the count distribution

start by visualizing distribution (here we're using the wild type data)


ggplot(raw_counts) +

geom_histogram(aes(x = wt_normal1), stat = 'bin', bins = 200) +

xlab("Raw expression counts") +

ylab("Number of genes")

output>



noted long right tail (this is due to there being no limit for maximum expression in RNA-Seq data)

looks like a Poisson distribution

however since there is expression variation between biological replicates Poisson would not be an appropriate model

*here the best model to use is the negative binomial model

Prep for DESeq2 object

```
dds <- DESeqDataSetFromMatrix(countData = rawcounts,
                              colData = metadata,
                              design = ~condition)
```

*the design formula given should contain major expected sources of variation to control for and the condition of interest as the last term in the formula

The metadata

```
# Create vectors containing metadata for the samples
genotype <- c("wt", "wt", "wt", "wt", "wt", "wt", "wt")
condition <- c("normal", "fibrosis", "normal",
               "fibrosis", "normal", "fibrosis", "fibrosis")

# Combine vectors into a data frame
wt_metadata <- data.frame(genotype, wildtype)
```

at the very least we need to know which of our samples correspond to each condition

sample names are added as the row names

Example

Create genotype vector

```
genotype <- c("smoc2_oe", "smoc2_oe", "smoc2_oe", "smoc2_oe", "smoc2_oe",  
"smoc2_oe", "smoc2_oe")
```

Create condition vector

```
condition <- c("fibrosis", "fibrosis", "fibrosis", "fibrosis", "normal", "normal",  
"normal")
```

Create data frame

```
smoc2_metadata <- data.frame(genotype, condition)
```

Assign the row names of the data frame

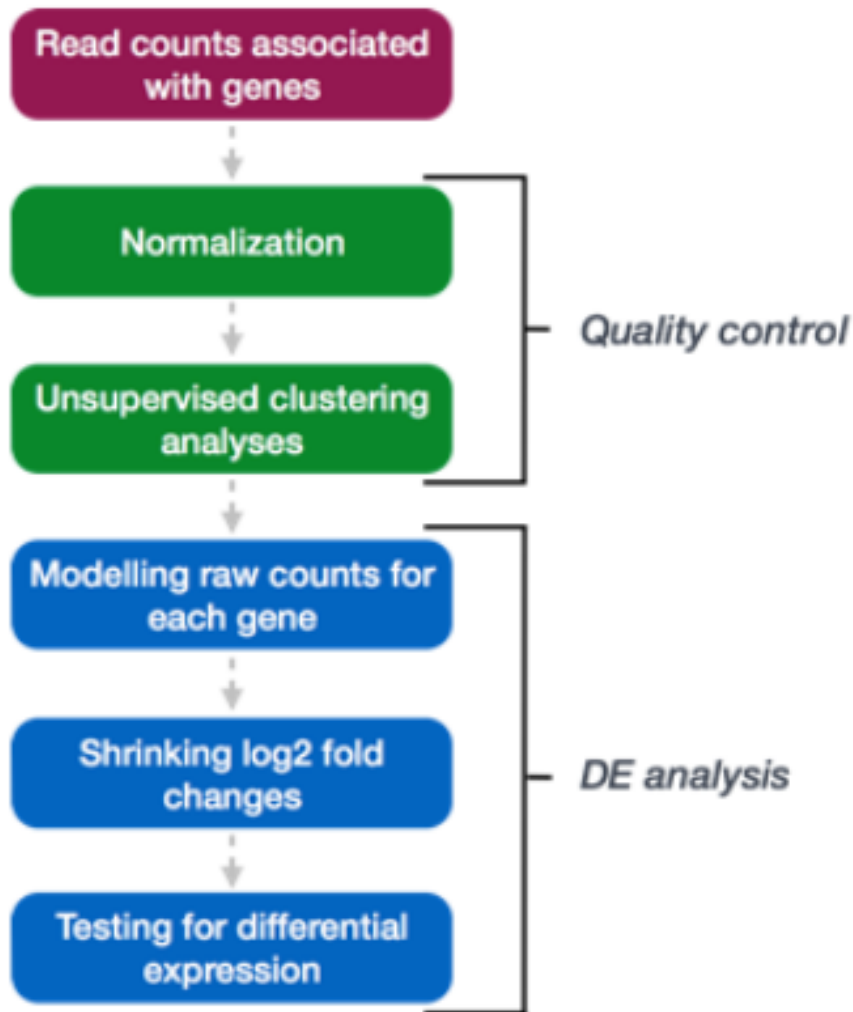
```
rownames(smoc2_metadata) <- c("smoc2_fibrosis1", "smoc2_fibrosis2",  
"smoc2_fibrosis3", "smoc2_fibrosis4", "smoc2_normal1", "smoc2_normal3",  
"smoc2_normal4")
```

Using DESeq2

a great function to get detailed info is;

```
vignette(DESeq2)
```

Workflow



first we normalize counts to account for differences in library depth
then we use PCA and hierarchical clustering using heatmaps to identify potential sample outliers and major variations in the data
then the DE analysis

Bringing data in
metadata and counts datasets need to be in the same order
row names of our metadata need to be in the same order as the column names of our counts data
can explore this data easily with these functions:
`rownames(wt_metadata)`
`colnames(wt_rawcounts)`
easy way to check if our rows and columns are in proper order >
`all(rownames(wt_meatdata) == colnames(wt_rawcounts))`
output willl be a logical response

How to match them?

```
match(vector1, vector2)
```

#vector1 represents the set with the desired order

#vector2 represents the set of values to reorder

output > the indices for how to rearrange vector2 to be in the same order as vector1

now we can reorder

example

```
idx <- match(colnames(wt_rawcounts), rownames(wt_metadata))
```

```
reordered_wt_metadata <- wt_metadata[idx, ]
```

#confirm in proper order

```
View(reordered_wt_metadata)
```

and

```
all(rownames(wt_metadata) == colnames(wt_rawcounts))
```

output shall now be TRUE

Create the DESeq2 object

```
dds_wt <- DESeqDataSetFromMatrix(countData = wt_rawcounts,  
                                  colData = reordered_wt_metadata,  
                                  design = ~ condition)
```

this object has class Ranged Summarized Experiment

this is a list-like object with slots available for the data it will generate throughout the analysis

Example

Use the match() function to reorder the columns of the raw counts

```
reorder_idx <- match(rownames(smoc2_metadata), colnames(smoc2_rawcounts))
```

Reorder the columns of the count data

```
reordered_smoc2_rawcounts <- smoc2_rawcounts[, reorder_idx]
```

Create a DESeq2 object

```
dds_smoc2 <- DESeqDataSetFromMatrix(countData =  
reordered_smoc2_rawcounts,  
                                     colData = smoc2_metadata,  
                                     design = ~ condition)
```

Count normalization

raw counts represent the number of reads aligning to each gene and should be proportional to the expression of the RNA in the sample

factors other than the expression of the RNA can influence the number of reads aligning to each gene

normalization methods can remove the influence of these factors

factors considered are library depth, gene length, and RNA composition

*DE analysis compares expression levels of the same genes between conditions, we do not need to normalize for gene length

*however to compare the expression levels of different genes, you would need to account for lengths of the genes

*also be aware that a few highly differentially expressed genes can skew many normalization methods

example gene DE



to get around this DESeq2 uses a 'median of ratios' method of normalization
this method adjusts the raw counts for library size and is resistant to large numbers of differentially expressed genes

To calculate

```
dds_wt <- estimateSizeFactors(dds_wt)
```

*this output gets assigned to a slot in the DESeq2 object dds_wt

the raw counts for each sample are divided by the associated sample-specific size factor for normalization

```
sizeFactors(dds_wt) #lets us view the size factors used for normalization
```

Now we can extract

we do this with the counts() function

```
normalized_wt_counts <- counts(dds_wt, normalized=TRUE)
```

ensure to put 'normalized' to true or else you will get the raw counts

Hierarchical heatmap

a solid tool to explore how similar the samples are to each other

first we need to log transform the normalized counts to improve the visualization of the clustering


DESeq2 uses a variance stabilizing transformation

this moderates the variance across the mean

'blind' argument specifies that the transformation should be blind to the sample information given in the design formula

```
vsd_wt <- vst(dds_wt, blind=TRUE)
```

Heatmap example



assesses the similarity in gene expression between the different samples in a dataset

used to explore how similar replicates are to each other and whether the samples belonging to different sample groups cluster separately

the heatmap is created by using the gene expression correlation values for all pairwise combinations of samples in the dataset

with 1 being a perfect correlation

*the hierarchical tree on the fringes shows which samples are more similar to each other

*we expect the biological replicates to cluster together and sample conditions to cluster apart

majority of genes should not be differently expressed, so samples should generally have high correlations with each other

*samples with correlation values below 0.8 require further investigation for contamination or outliers

to create a hierarchical heatmap

#extract the vst-normalized counts as a matrix from the vsd object

```
vsd_mat_wt <- assay(vsd_wt)
```

#then compute the pairwise correlation values between each pair of samples using the cor() function

```
vsd_cor_wt <- cor(vsd_mat_wt)
```

```
View(vsd_cor_wt)
```

output>

	wt_normal1	wt_normal2	wt_normal3	wt_fibrosis1	wt_fibrosis2	wt_fibrosis3	wt_fibrosis4
wt_normal1	1.0000000	0.9934287	0.9945298	0.9616998	0.9708459	0.9626185	0.9696097
wt_normal2	0.9934287	1.0000000	0.9930148	0.9629644	0.9713154	0.9639685	0.9704541
wt_normal3	0.9945298	0.9930148	1.0000000	0.9678018	0.9758950	0.9683519	0.9750891
wt_fibrosis1	0.9616998	0.9629644	0.9678018	1.0000000	0.9930090	0.9939055	0.9926560
wt_fibrosis2	0.9708459	0.9713154	0.9758950	0.9930090	1.0000000	0.9931793	0.9939010

Next


```
library(pheatmap)
```

'annotation' argument selects which factors in the metadata to include as annotation bars

use the select() function from the dplyr package to select the condition column in the wildtype metadata

```
pheatmap(vsd_cor_wt, annotation = select(wt_metadata, condition))
```

output for our example>



Example

```
# Transform the normalized counts
```

```
vsd_smoc2 <- vst(dds_smoc2, blind = TRUE)
```

```
# Extract the matrix of transformed counts
```

```
vsd_mat_smoc2 <- assay(vsd_smoc2)
```


```
# Compute the correlation values between samples
```

```
vsd_cor_smoc2 <- cor(vsd_mat_smoc2)
```

```
# Plot the heatmap
```

```
pheatmap(vsd_cor_smoc2, annotation = select(smoc2_metadata, condition))
```

output>




PCA


a technique used to emphasize the variation present in a dataset

find the principal components of a dataset

with PC1 (the first principal component) representing the greatest amount of variance in the data




Simple example where we have two samples and we plot the occurrence of the particular gene using the x-axis for 1 sample and the y-axis for the other sample




	Sample 1	Sample 2
Gene A	4	5
Gene B	1	4
Gene C	8	8
Gene D	5	7

Sample 1

Next we draw a line where there is the most variation (PC1)



the second most variation (PC2) must be perpendicular to PC1
this is to best describe the variance within the dataset
this is a very simple example, most datasets will be multi-dimensional
*the number of principal components is equal to the number of samples (n) in the dataset
the most variant genes for a principal component have the most influence on that principal component's direction
quantitative scores to genes based on how much they influence the different PCs



	Sample 1	Sample 2	Influence on PC1	Influence on PC2
Gene A	4	5	-2	0.5
Gene B	1	4	-10	1
Gene C	8	8	8	-5
Gene D	5	7	1	6

per sample PC value is computed by taking the product of the influence and the normalized read count for each gene and summing across all genes

```

Sample1 PC1 score = (4 * -2) + (1 * -10) + (8 * 8) + (5 * 1) = 51
Sample1 PC2 score = (4 * 0.5) + (1 * 1) + (8 * -5) + (5 * 6) = -7

Sample2 PC1 score = (5 * -2) + (4 * -10) + (8 * 8) + (7 * 1) = 21
Sample2 PC2 score = (5 * 0.5) + (4 * 1) + (8 * -5) + (7 * 6) = 8.5

```

we then plot these per sample PC values
samples that cluster together have more similar gene expression profiles than
samples that cluster apart


PCA can help identify sample outliers and major sources of variation

Using DESeq2

#plotPCA() takes as an input the transformed vsd object

#'intgroup' argument specifies what factor in the metadata to use to color the plot

```
# Plot PCA
plotPCA(vsd_wt, intgroup="condition")
```



what this shows

on PC1 we can see that the sample groups fibrosis and normal separate nicely
this means that our condition corresponds to PC1 representing 88% of the variance in our data

4% of the variance is explained in PC2

basically saying the majority of the variation in gene expression in the dataset can likely be explained by the differences between sample groups

if they didn't separate that could mean that the effect of the condition could be small

or other larger sources of variation may be present

can color other factors (ie age, sex,...) to identify these

also function `pcomp()` allows a more thorough analysis of PCs other than PC1 and PC2

DE analysis

reminder first steps (from above)

explore quality of samples

remove any outlier samples

assess the major sources of variation present

DE analysis steps

1. fitting the raw counts for each gene to the DESeq2 negative binomial model and testing for differential expression
2. shrinking the log2 fold changes
3. extracting and visualizing the results

Design formula

tells DESeq2 the known major sources of variation to control for (or regress out) and the condition of interest to use for differential expression testing

*condition of interest comes last in our design formula

example from above

```
dds_wt <- DESeqDataSetFromMatrix(countData = wt_rawcounts,  
                                  colData = reordered_wt_metadata,  
                                  design = ~ strain + sex + treatment)
```

#here treatment is our condition of interest

#the other factors are placed because in our prior assessment they proved to be major sources of variation

#these other factors can be placed in any order but the condition of interest needs to be last

tilde tells DESeq2 to model the counts using the following formula

factor names in the design formula need to exactly match the column names in the metadata

We can do more complex designs with DESeq2

example say we wanted to know the effect of sex on the effect of treatment we could use an interaction term

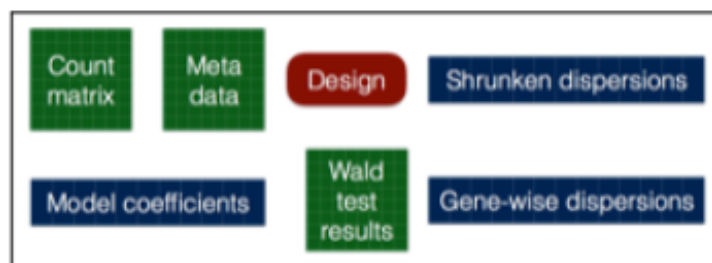
~ strain + sex + treatment + sex:treatment

Now we can perform model fitting


example

```
# Run analysis  
dds_wt <- DESeq(dds_wt)
```

```
using pre-existing size factors  
estimating dispersions  
gene-wise dispersion estimates  
mean-dispersion relationship  
final dispersion estimates  
fitting model and testing
```



Assessing how well the model fits



the goal of the differential expression analysis is to determine whether a gene's mean expression between sample groups is different given the variation within groups

determined by testing the probability of the log2 fold changes between groups being significantly different from zero

to model the counts requires information about the mean and variation in the data to explore this variation we observe the variance in gene expression relative to the mean

variance (square of std) represents how far away the expression of the individual samples are from the means

For RNA-Seq data

variance is generally expected to increase with the gene's mean expression

the `apply()` function allows us to calculate the means and variances for every gene of the normal samples

```
# Syntax for apply()
```

```
apply(data, rows/columns, function_to_apply)
```

```
# Calculating mean for each gene (each row)
```

```
mean_counts <- apply(wt_rawcounts[, 1:3], 1, mean)
```

```
# Calculating variance for each gene (each row)
```

```
variance_counts <- apply(wt_rawcounts[, 1:3], 1, var)
```


then create a dataframe

then plot

```
# Creating data frame with mean and variance for every gene  
df <- data.frame(mean_counts, variance_counts)
```

```
ggplot(df) +  
  geom_point(aes(x=mean_counts, y=variance_counts)) +  
  scale_y_log10() +  
  scale_x_log10() +  
  xlab("Mean counts per gene") +  
  ylab("Variance per gene")
```

a log10 scale helps in visualizing this range more effectively
each black dot represents a gene



the variance in gene expression increase with the mean
this is expected in RNA-Seq data
the range in values for variance is greater for lower mean counts compared to
higher mean counts
also expected
'dispersion' is a measure of the variance for a given mean

DESeq2 uses dispersion to assess the variability in expression when modeling the counts

Var: variance

μ : mean

α : dispersion

Dispersion formula: $Var = \mu + \alpha * \mu^2$

Relationship between mean, variance and dispersion:


$\uparrow \text{variance} \Rightarrow \uparrow \text{dispersion}$

$\uparrow \text{mean} \Rightarrow \downarrow \text{dispersion}$

for any two genes with the same mean expression, the only difference in dispersion will be based on differences in variance

*to check fit of our data to our DESeq2 model, it can be useful to look at the dispersion estimates

```
# Plot dispersion estimates  
plotDispEsts(dds_wt)
```



each black dot is a gene with associated mean and dispersion values
we expect dispersion values to decrease with increasing mean, which is what we

see

few replicates can mean gene-wise estimates of dispersion are often inaccurate
DESeq2 uses information across all genes to determine the most likely estimates of dispersion #red line

*genes with inaccurately small estimates of variation could yield many false positives or genes that are identified as DE (when they are not)

the black dots are the originals

the blue dots are the originals shrunk towards the curve

this yields more accurate estimates of dispersion

the blue dot estimates of dispersion are used to model the counts for determining the differentially-expressed genes

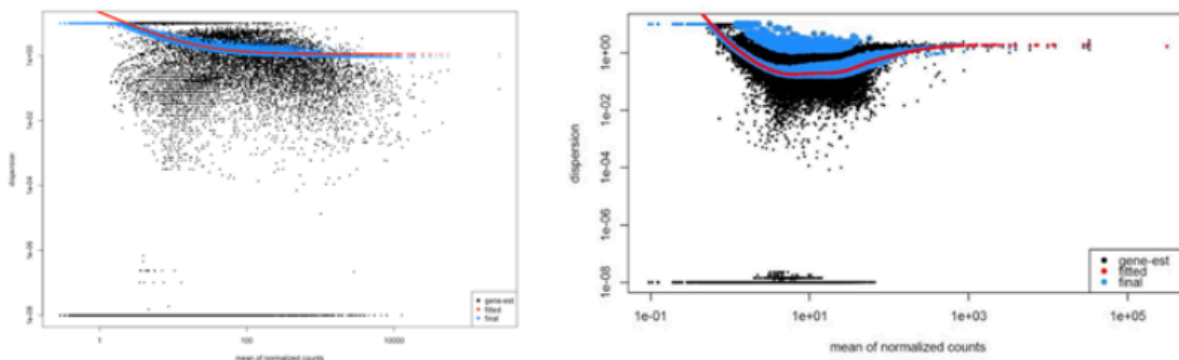
*extremely high dispersion values (blue circles with black dots inside) are not shrunk

this is due to the likelihood the the gene may have higher variability than others (biological or technical reasons)

*and reducing this variation could result in false positives

larger number of replicates can estimate the mean and variation more accurately
so yield less shrinkage

Worrisome dispersion plots



the first does not follow the curve

the second dispersion does not decrease with increasing mean

DESeq2 Negative Binomial Model

raw count for gene i , sample j

The mean is taken as "normalized counts" scaled by a normalization factor

one dispersion per gene

$$K_{ij} \sim \text{NB}(s_{ij}q_{ij}, \alpha_i)$$

negative binomial model accounts for the additional variation in the data added by the small number of biological replicates
 size factors indicated by s_{ij}
 normalized counts indicated by q_{ij}
 shrunk dispersions denoted by α_i
 all used as input to the negative binomial model to fit the raw count data

normalized counts for gene i , sample j

\log_2 fold change between conditions

$$\log_2 q_{ij} = \sum_r X_{jr} \beta_{ir}$$

for each gene the model uses the \log_2 normalized counts to determine the \log_2 foldchange estimates (β_{ir}) for the samples of the condition of interest (X_{jr})
 in addition the associated standard error is also the output

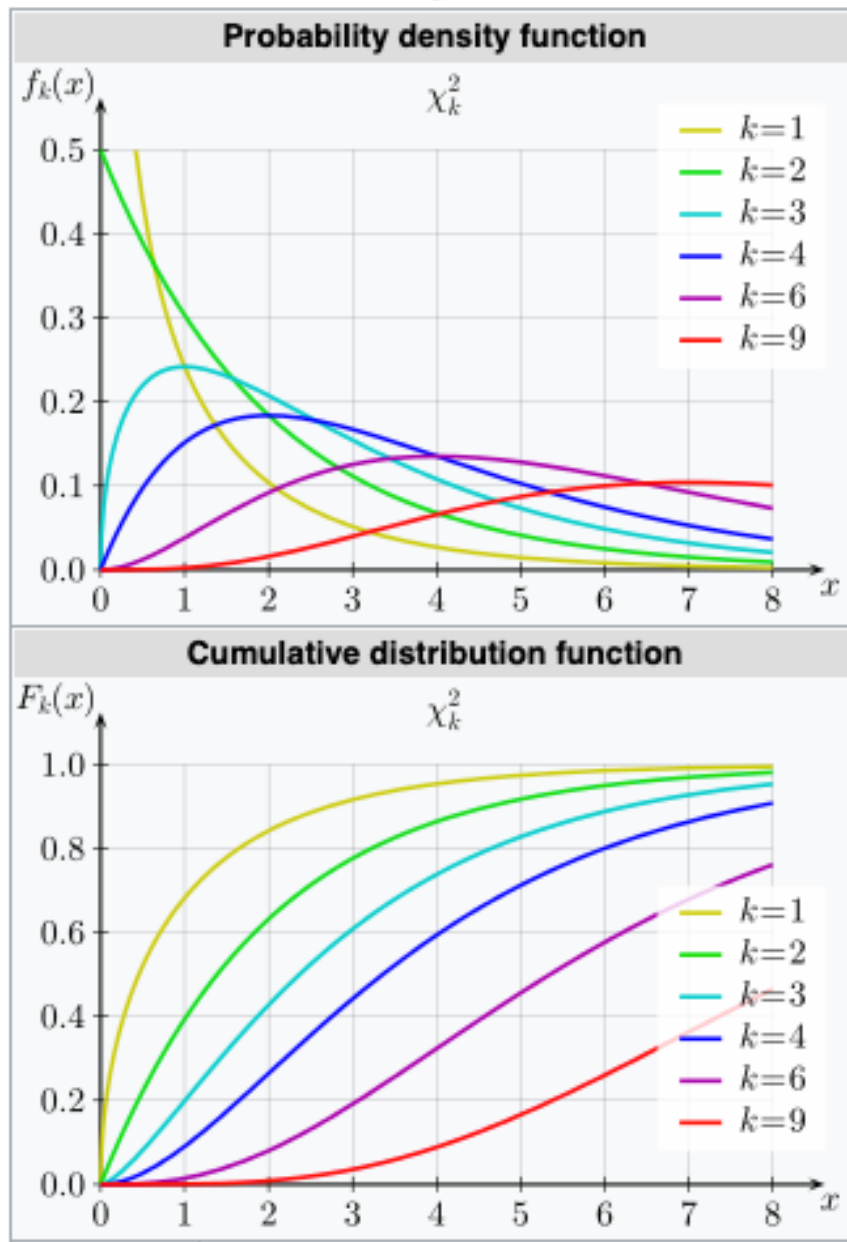
DESeq2 contrasts

by default DESeq2 performs the Wald test for pairwise comparisons to test for differences in expression between two sample groups for the condition of interest
 the Wald test is a statistical test used to assess the significance of individual coefficients (parameters) in a statistical model
 particularly in the context of regression analysis

$$\text{Wald statistic} = \left(\frac{\text{Estimated Coefficient} - \text{Hypothesized Value}}{\text{Standard Error of the Coefficient}} \right)^2$$

with large sample sizes Wald statistic follows a chi-squared distribution
 chi-squared distribution is a probability distribution used in hypothesis testing and confidence interval construction
 looks like:

Chi-squared



chi-squared distribution is positively skewed and takes only non-negative values characterized by degrees of freedom (denoted as k or v)

this determines the shape of the distribution

if assumed standard normal random variables Z (mean 0 and variance 1), the random variable X follows a chi-squared distribution

as the dofs increase the chi-squared distribution approaches a normal distribution

CDF is used to calculate the prob that a chi-squared random variable is less than or equal to a specified value

Within DESeq2 to get the results of this testing:

`results(wt_dds, alpha = 0.05)`

alpha is chosen based on how stringent you want to be with your analysis
lower alpha indicate less probability of identifying a gene as DE when it is actually not

output example with condtion meaning normal vs fibrosis>

```
log2 fold change (MLE): condition normal vs fibrosis
Wald test p-value: condition normal vs fibrosis
DataFrame with 47729 rows and 6 columns
```

	baseMean	log2FoldChange	lfcSE
	<numeric>	<numeric>	<numeric>
ENSMUSG000000102693	0	NA	NA
ENSMUSG000000064842	0	NA	NA
ENSMUSG000000051951	19.5084656230804	-4.1291703663434	0.815892075046039
ENSMUSG000000102851	0	NA	NA
ENSMUSG000000103377	0	NA	NA

this represents the normal group relative to the fibrosis group

We can perform any pairwise comparison between the sample groups by supplying our own contrast

```
results(dds,  
        contrast = c("condition_factor", "level_to_compare",  
                      "base_level"),  
        alpha = 0.05)
```

```
wt_res <- results(dds_wt,  
                 contrast = c("condition", "fibrosis",  
                              "normal"),  
                 alpha = 0.05)
```


here we define 'normal' as the base level and 'fibrosis' as the level to compare

Exploring these results a bit further

MA plot shows the mean of the normalized counts versus the log2 fold changes for all genes tested

```
plotMA(wt_res, ylim=c(-8,8))
```

output>



genes that are significantly DE are colored red

large log2 foldchanges for genes with lower mean count values

*these fold changes are unlikely to be as accurate for genes that have little information associated with them

such as genes with low numbers of counts or high dispersion values

refresher on fold changes:

represent the ratio of one value to another

tells you how many times one quantity is larger or smaller than another

Fold Change = B/A

basically says "how many times is B different than A?"

if 1, then quantities are the same

if 2, then B is twice as large as A

if <1 , then B is smaller than A


LFC shrinkage

to improve the estimated fold changes

for genes with low amounts of information available, shrinkage use info from all genes to generate more likley (lower) log2 fold changes

```
wt_res <- lfcShrink(dds_wt,
  contrast=c("condition", "fibrosis", "normal"),
  res=wt_res)
```

now using the plot MA



we can see more restricted log2 fold change values particularly for lowly expressed genes
 these fold changes should be more accurate
 *this will not affect the number of differentially expressed genes returned
 now we can extract the significant DE genes and perform further visualization of results

Example

```
# Shrink the log2 fold change estimates to be more accurate
smoc2_res <- lfcShrink(dds_smoc2,
  contrast = c("condition", "fibrosis", "normal"),
  res = smoc2_res)
```

DESeq2 results table

to get descriptions for the columns in the results table

```
mccols(wt_res)
```

*to determine significant DE genes, we use the p-values adjusted for multiple test correction in the last column

```
head(wt_res, n=10)
```

```
output>
```

```
log2 fold change (MAP): condition fibrosis vs normal
Wald test p-value: condition fibrosis vs normal
data frame with 6 rows and 6 columns
```

	baseMean	log2FoldChange	lfcSE	stat
	<numeric>	<numeric>	<numeric>	<numeric>
ENSMUSG000000102693	0	NA	NA	NA
ENSMUSG000000064842	0	NA	NA	NA
ENSMUSG000000051951	19.5084656230804	3.55089043143673	0.648400500074659	4.66871842838828
ENSMUSG000000102851	0	NA	NA	NA
ENSMUSG000000103377	0	NA	NA	NA
ENSMUSG000000104017	0	NA	NA	NA

	pvalue	padj
	<numeric>	<numeric>
ENSMUSG000000102693	NA	NA
ENSMUSG000000064842	NA	NA
ENSMUSG000000051951	3.03084428526558e-06	1.93776447202312e-05
ENSMUSG000000102851	NA	NA
ENSMUSG000000103377	NA	NA
ENSMUSG000000104017	NA	NA

*remember what alpha 0.05 means, that for every gene tested has a 5% chance that the gene is called DE with it is not
when this happens, we have a false positive
put this into perspective, if we tested 50,000 genes we could have about 2,000 genes (5%) as false positives
because of this DESeq2 use multiple tes correction
uses the Benjamini-Hochberg method
this helps control the proportion of false positives relative to true positives
DESeq2 automatically filters out genes unlikely to be truly differentially expressed prior to testing

Which ones?:

- genes with 0 counts across all samples
- genes with low mean values across all samples
- genes with extreme count outliers

*above table, we can denote the filtered out genes by an 'NA'

summary() function provides the number of DE genes for our alpha and info about the number of genes filtered

```
out of 29866 with nonzero total read count
adjusted p-value < 0.05
LFC > 0 (up)      : 5413, 18%
LFC < 0 (down)    : 5235, 18%
outliers [1]      : 47, 0.16%
low counts [2]    : 8412, 28%
(mean count < 2)
[1] see 'cooksCutoff' argument of ?results
[2] see 'independentFiltering' argument of ?results
```

our results give over 10,000 genes as DE
this is the sum of the DE genes with log2 fold changes less than 0 and greater than 0

This is a lot of genes
not ideal but sometimes necessary we can add a threshold

```
wt_res <- results(dds_wt,
                  contrast = c("condition", "fibrosis", "normal"),
                  alpha = 0.05,
                  lfcThreshold = 0.32)
```

```
wt_res <- lfcShrink(dds_wt,
                   contrast=c("condition", "fibrosis", "normal"),
                   res=wt_res)
```

here we use a 1.25 fold change threshold which equals 0.32 on the log2 scale
this increases the risk of losing biologically relevant genes
but by keeping the threshold low we hope to reduce the risk

To better understand which genes the results pertain to we use the annotables package

```
library(annotables)
```

example

```
grcm38 #known mouse genome
```



```
# A tibble: 53,728 x 9
  ensgene      entrez symbol chr      start      end strand biotype      description
  <chr>      <int> <chr> <chr>      <int>      <int> <int> <chr>      <chr>
1 ENSMUSG00000000001 14679 Gnai3 3      108107280 108146146 -1 protein_coding guanine nucleotide binding protein (G p...
2 ENSMUSG00000000003 54192 Pbsn  X      77837901 77853623 -1 protein_coding probasin [Source:MGI Symbol;Acc:MGI:186...
3 ENSMUSG00000000028 12544 Cdc45 16     18780447 18811987 -1 protein_coding cell division cycle 45 [Source:MGI Symb...
4 ENSMUSG00000000031 NA H19 7      142575529 142578143 -1 lincRNA H19, imprinted maternally expressed tra...
5 ENSMUSG00000000037 107815 Scml2 X      161117193 161258213 1 protein_coding sex comb on midleg-like 2 (Drosophila) ...
6 ENSMUSG00000000049 11818 Apoh 11     108343354 108414396 1 protein_coding apolipoprotein H [Source:MGI Symbol;Acc...
7 ENSMUSG00000000056 67608 Narf 11     121237253 121255856 1 protein_coding nuclear prelamin A recognition factor [...]
```

Next, extracting results

```
wt_res_all <- data.frame(wt_res) %>%
  rownames_to_column(var = "ensgene") %>%
  left_join(x = wt_res_all,
            y = grcm38[, c("ensgene", "symbol", "description")],
            by = "ensgene")
View(wt_res_all)
```

this annotates the genes with the gene names and descriptions

turn to a df

change the row names to a column

then merge the gene names and descriptions with our results using a left join and merge by ensemble gene IDs (ensgene)

Now extract the significant DE genes

subset our results table for genes with p-adjusted values less than 0.05

```
wt_res_sig <- subset(wt_res_all, padj < 0.05)
wt_res_sig <- wt_res_sig %>%
  arrange(padj)
View(wt_res_all)
```

use arrange() to order the genes by p-adjusted values

ensgene	baseMean	log2FoldChange	lfcSE	stat	pvalue	padj	symbol	description
ENSMUSG00000053113	1318.1717	4.875042	0.16021506	28.35016	8.330958e-177	1.830145e-172	Socs3	suppressor of cytokine signaling 3 [Source:M...
ENSMUSG00000005087	2943.7403	6.121134	0.20721978	27.89891	2.750356e-171	3.020991e-167	Cd44	CD44 antigen [Source:MGI Symbol;Acc:MGI:8...
ENSMUSG00000036887	3899.5135	3.866162	0.12740248	27.83465	1.652344e-170	1.209957e-166	C1qa	complement component 1, q subcomponent,...
ENSMUSG00000026822	8870.1712	6.466148	0.23782361	25.82294	4.901029e-147	2.691645e-143	Lcn2	lipocalin 2 [Source:MGI Symbol;Acc:MGI:96757]
ENSMUSG00000036905	3237.6046	3.835279	0.13773926	25.52164	1.134018e-143	4.982421e-140	C1qb	complement component 1, q subcomponent,...
ENSMUSG00000027962	9298.5984	5.781446	0.21949603	24.88019	1.219253e-136	4.463724e-133	Vcam1	vascular cell adhesion molecule 1 [Source:MG...
ENSMUSG00000018008	1278.6520	3.202855	0.11631046	24.77939	1.495690e-135	4.693902e-132	Cytn4	cytohesin 4 [Source:MGI Symbol;Acc:MGI:244...
ENSMUSG00000051439	4144.4097	3.743987	0.14014630	24.43589	7.109040e-132	1.952142e-128	Cd14	CD14 antigen [Source:MGI Symbol;Acc:MGI:8...

we should also see log2 foldchanges are >0.32

now we can explore this table for interesting or expected genes with a high

probability of being related to kidney fibrosis
lit review time

Example

```
# Explore the results() function  
?results
```

```
# Extract results
```

```
smoc2_res <- results(dds_smoc2,  
  contrast = c("condition", "fibrosis", "normal"),  
  alpha = 0.05,  
  lfcThreshold = 0.32)
```

```
# Shrink the log2 fold changes
```

```
smoc2_res <- lfcShrink(dds_smoc2,  
  contrast = c("condition", "fibrosis", "normal"),  
  res = smoc2_res)
```

```
# Save results as a data frame
```

```
smoc2_res_all <- data.frame(smoc2_res)
```

```
# Subset the results to only return the significant genes with p-adjusted values  
less than 0.05
```

```
smoc2_res_sig <- subset(smoc2_res_all, padj < 0.05)
```

Visualization of results

start with expression heatmap

explores the expression of significant genes

here we plot the normalized count values of the significant genes instead of the
sample correlation values

subset the normalized counts to only the significant DE genes

example

```
sig_norm_counts_wt <- normalized_counts_wt[wt_res_sig$ensgene, ]
```

```
#we can choose a color palette with RColorBrewer
```

```
library(RColorBrewer)
```


```
heat_colors <- brewer.pal(6, "YlOrRd")
```

```
#to see all the possible color palettes
```

```
display.brewer.all()
```


```
# Run pheatmap
pheatmap(sig_norm_counts_wt,
         color = heat_colors,
         cluster_rows = T,
         show_rownames = F,
         annotation = select(wt_metadata, condition),
         scale = "row")
```

cluster by row, annotating with condition information and scaling by row
 this plots Z-scores rather than the actual normalized count values
 we expect to see the expression levels for the significant genes to cluster by
 sample group
 output>



Additional useful plot - Volcano plot
 shows the fold changes relative to the adjusted p-values for all genes

```
ggplot(wt_res_all) +
  geom_point(aes(x = log2FoldChange, y = -log10(padj)), color = threshold)) +
  xlab("log2 fold change") +
  ylab("-log10 adjusted p-value") +
  ylim=c(0, 15) +
  theme(legend.position = "none",
        plot.title = element_text(size = rel(1.5), hjust = 0.5),
        axis.title = element_text(size = rel(1.25)))
```



ylim() allows us to better visualize the significance cut-off

Visualizing the top significant genes

to do this create a dataframe

our example we want the top 20 genes

```
top_20 <- data.frame(sig_norm_counts_wt)[1:20, ] %>%
```

```
  rownames_to_column(var = 'ensgene')
```

```
top_20 <- gather(top_20, key = "samplename", value = "normalized_counts", 2:8)
```

#2:8 is our last argument and specifies the columns we want to gather into a single column of values

to plot the expression we want to merge the metadata so that we can color the plot by sample group

we can use the inner_join() function to keep only those columns in both datasets we join on the row names

we need to turn these into a column to join on them

we put the normalized counts on the y-axis

gene IDs on the x-axis

```
top_20 <- inner_join(top_20,
```

```
  rownames_to_column(wt_metadata, var = 'samplename'),
```

```
  by = "samplename")
```

```
ggplot(top_20) +
```

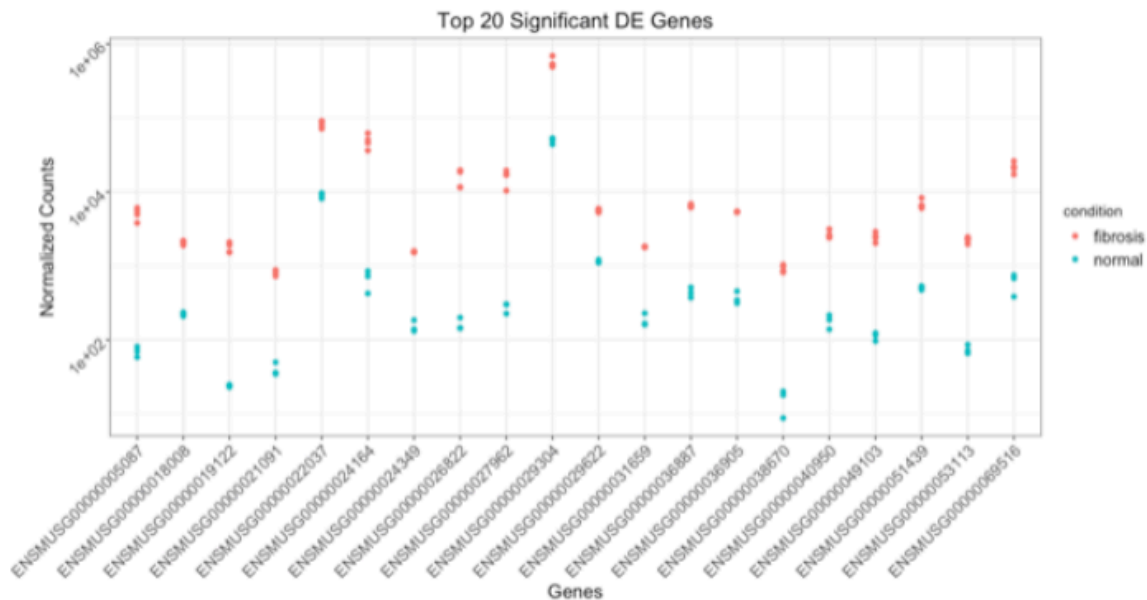
```
  geom_point(aes(x = ensgene, y = normalized_counts, color = condition)) +
```

```

scale_y_log10() +
xlab("Genes") +
ylab("Normalized Counts") +
ggtitle("Top 20 Significant DE Genes") +
theme_bw() +
theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
theme(plot.title = element_text(hjust = 0.5))

```

#we use the log10 scale on the y-axis to more easily visualize the wide range in expression values



Example

```

# Subset normalized counts to significant genes
sig_norm_counts_smoc2 <-
normalized_counts_smoc2[rownames(smoc2_res_sig), ]

```

```

# Choose heatmap color palette
heat_colors <- brewer.pal(n = 6, name = "YlOrRd")


```

```

# Plot heatmap
pheatmap(sig_norm_counts_smoc2,
  color = heat_colors,
  cluster_rows = TRUE,
  show_rownames = FALSE,
  annotation = select(smoc2_metadata, condition),
  scale = "row")

```

output>



DE Analysis Summary

From alignment or mapping

alignment or mapping of the reads to the genome to determine the location on the genome where the reads originated

reads derived from mRNA are often aligned to the organism's genome

some of these reads cross introns

tools for aligning reads to the genome need to align across introns or be splice-aware for RNA-seq

the output of alignment gives the genome coordinates for where the read most likely originated from in the genome and info about the quality of the mapping

*the reads aligning to the exons of each gene are quantified to yield a matrix of gene counts

*the number of reads aligning to each of the genes is given in the count matrix and represents the expression level of the gene

*the more reads aligning to the gene, the higher the expression of the gene meaning more RNA transcripts were expressed

once we have the counts, differential expression analysis is performed

goal is to determine whether the gene counts between the sample groups are significantly different given the variation in the counts within the sample group

Example

Check that all of the samples are in the same order in the metadata and count data

```
all(rownames(all_metadata) %in% colnames(all_rawcounts))
```

```
# DESeq object to test for the effect of fibrosis regardless of genotype
```

```
dds_all <- DESeqDataSetFromMatrix(countData = all_rawcounts,  
                                  colData = all_metadata,  
                                  design = ~ condition)
```

```
# DESeq object to test for the effect of genotype on the effect of fibrosis
```

```
dds_complex <- DESeqDataSetFromMatrix(countData = all_rawcounts,  
                                       colData = all_metadata,  
                                       design = ~ genotype + condition + genotype:condition)
```

Recap of workflow - key notes

Normalization allows comparison of counts within a sample and between the different samples

to compare counts for the same gene between different samples we need to normalize for library size, or total number of reads per sample

the normalized counts are needed to explore the similarity between samples using hierarchical correlation heatmaps and PCA

we log transform the normalized counts to improve the distances of the clusters for visualization

vst() function performs this transformation while being blind to sample group information

for the heatmap, we extract the transformed normalized counts using the assay() function on the transformed counts object

then compute the correlation values between samples using the cor() function

```
vsd %>%  
  assay() %>% # Extract the vst matrix from the object  
  cor() %>%   # Compute pairwise correlation values  
  pheatmap(annotation = metadata[, c("column_name1", "column_name2")])
```

*remember we are looking for our biological replicates to cluster together and the sample groups to cluster apart

also identifies outliers which tend to have low correlation values with all other samples

Example

```
# Log transform counts for QC
```

```
vsd_all <- vst(dds_all, blind = TRUE)
```

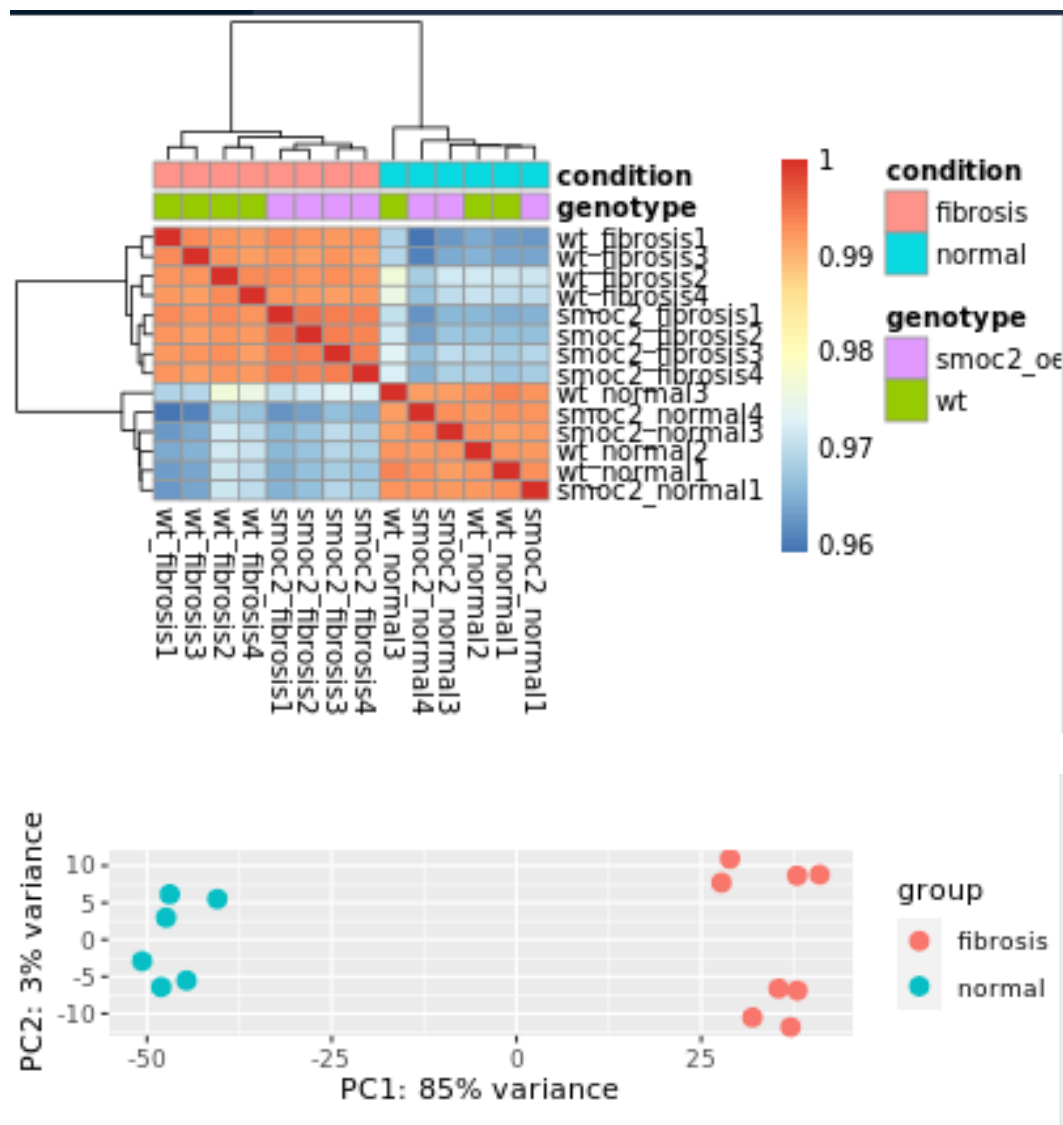
```
# Create heatmap of sample correlation values
```


```
vsd_all %>%
  assay() %>%
  cor() %>%
  pheatmap(annotation = select(all_metadata, c("genotype", "condition")))
```

```
# Create the PCA plot for PC1 and PC2 and color by condition
plotPCA(vsd_all, intgroup = "condition")
```

```
# Create the PCA plot for PC1 and PC2 and color by genotype
plotPCA(vsd_all, intgroup = "genotype")
```

output>





```
# Select significant genes with padj < 0.05
smoc2_sig <- subset(res_all, padj < 0.05) %>%
  data.frame() %>%
  rownames_to_column(var = "geneID")
```

```
# Extract the top 6 genes with padj values
smoc2_sig %>%
  arrange(padj) %>%
  select(geneID, padj) %>%
  head(6)
```